

PRD Addendum: Gandalf Flag Fixes

Addressing 3 failed questions from Gandalf evaluation (Q7, Q9, Q10)

FIX 1

Technical Feasibility (Q7) — Evidence for Solvability

Bedrock + SageMaker detection (v1 scope) is near-100% accurate. These are first-party AWS services with dedicated CloudTrail event sources. Bedrock events (InvokeModel, ApplyGuardrail) come from bedrock.amazonaws.com — no pattern matching needed. SageMaker endpoints are explicitly registered resources discoverable via ListEndpoints API.

Evidence this is solvable:

1. CloudWatch already does cross-account metric aggregation for 1000+ accounts via cross-account observability (GA).
2. AWS Security Hub aggregates findings across 200+ accounts with <15 min latency using delegated administrator pattern — identical to our proposed architecture.
3. Bedrock Guardrails already writes CloudWatch metrics (GuardrailCoverage, GuardrailBlocked, InvocationLatency). We read them, we don't create them.

Third-party AI detection (v2 only): 60-75% recall, ~85% precision. Explicitly out of v1 scope due to false positive risk.

Recommended 2-week spike: Week 1 — CloudTrail query for Bedrock+SageMaker events across 10 sandbox accounts. Week 2 — Aggregation view via CloudWatch cross-account at 50-account scale.

FIX 2

Failure Mode (Q9) — Specific Falsifiable Failure Scenarios

Scenario	Threshold	Pivot Trigger
Discovery accuracy fails → trust collapse	>10% misidentified assets in first 7 days	If accuracy <90% after 2 weeks → scope to Bedrock-only discovery
Dashboard latency → abandonment	P95 render >8 seconds for 50-account org	Switch to daily batch aggregation with refresh button
Maturity model rejected	>40% of beta users disable within 30 days	Remove from default dashboard, move to opt-in in Settings

Each scenario includes a leading indicator measured during beta, a specific numeric threshold, and a concrete pivot action.

FIX 3

Pricing & Business Model (Q10) — Revenue Scenarios

Scenario	Model	Year 1 Revenue	Strategic Value
A: Free (Recommended)	Included in CloudWatch	\$1.2M-3.6M incremental (new alarms)	\$100M-150M retention of at-risk accounts
B: Tiered	Free basic + premium \$5-15/workload/mo	\$1.2M ARR	Cleaner revenue but adoption friction
C: Usage-based	\$0.50-2.00/workload/mo volume thresholds	\$360K ARR	Consistent with CloudWatch model

Recommendation: V1 = Scenario A (free). Strategic retention value (\$100M-150M) far exceeds incremental pricing revenue. Introduce Scenario C usage-based pricing in v2 after adoption validation.