

# Gandalf Evaluation: CloudWatch AI Control Plane

Adversarial Strategy & Problem Gate — Stage 3

## VERDICT: PASSED WITH FLAGS

Score: 7/10 questions passed | Rounds used: 1/3

Three questions flagged for human review. Pipeline advances.

### Detailed Scores

| #  | Dimension                | Rubric (1-5) | Evidence (0/1) | Pass? | Key Notes                                                                                   |
|----|--------------------------|--------------|----------------|-------|---------------------------------------------------------------------------------------------|
| 1  | TAM & Market             | 3            | 1              | PASS  | SAM \$75M-300M shown but input estimates unvalidated                                        |
| 2  | Why Now                  | 4            | 1              | PASS  | 3 convergent 12-month signals. Strongest answer.                                            |
| 3  | Customer Problem Depth   | 3            | 1              | PASS  | Vivid persona, quantified cost. Scenarios are constructed, not observed. Barely passes.     |
| 4  | North Star Metric        | 4            | 1              | PASS  | Well-justified. Rejected alternatives named.                                                |
| 5  | Competitive Moat         | 4            | 1              | PASS  | Structural moat (AWS infra access), not just features.                                      |
| 6  | Scope Discipline         | 3            | 1              | PASS  | Explicit cuts with rationale. V1 may still be too big.                                      |
| 7  | Technical Feasibility    | 2            | 0              | FAIL  | Zero evidence CloudTrail AI detection works. No prototype, no spike results, no benchmarks. |
| 8  | Cannibalization Risk     | 4            | 1              | PASS  | Anti-metric tracking. Bedrock team alignment flagged.                                       |
| 9  | Failure Mode             | 2            | 0              | FAIL  | Generic risks without falsifiable thresholds. No specific "we pivot if X" criteria.         |
| 10 | Pricing & Business Model | 2            | 1              | FAIL  | "Included in CloudWatch pricing" — no revenue model. Retention value unquantified.          |

### Questions Flagged for Human Review

#### Q7: Technical Feasibility — FAIL (Rubric: 2, Evidence: 0)

**Current answer:** FAQ Q16 describes 'no new data store' architecture. Q21 names cross-account aggregation latency as biggest risk with '15-minute cached refresh' mitigation.

**What's missing:** Zero evidence that CloudTrail pattern detection for AI API calls works at the claimed accuracy. No prototype or spike results for cross-account AI discovery. No latency benchmarks for aggregating 100+ accounts. No reference implementations.

**Suggested action:** Run a 2-week technical spike: (1) Can CloudTrail reliably detect Bedrock InvokeModel and SageMaker InvokeEndpoint across 10+ accounts in < 5 minutes? (2) What's the false positive rate for third-party AI API detection? (3) Build a prototype aggregation query for 50 accounts and measure latency.

#### Q9: Failure Mode — FAIL (Rubric: 2, Evidence: 0)

**Current answer:** Six risks listed in a table with likelihood/impact/mitigation, but all are generic scenarios without falsifiable thresholds.

**What's missing:** Specific failure criteria with numbers. 'If AI discovery misidentifies > 15% of assets in a customer's first session, they won't return' is falsifiable. 'Discovery false positives erode trust' is not.

**Suggested action:** Define 3 specific failure scenarios with measurable thresholds and leading indicators that would trigger a pivot or scope reduction.

### ***Q10: Pricing & Business Model — FAIL (Rubric: 2, Evidence: 1)***

**Current answer:** 'Included in existing CloudWatch pricing' with acknowledgment that cost optimization could reduce AI service revenue.

**What's missing:** Revenue modeling. If this is a retention play, quantify the retention value (\$X in at-risk annual revenue from AI-heavy customers who might leave for Datadog/ServiceNow). If incremental revenue is possible, model the scenarios.

**Suggested action:** Work with Finance to model three scenarios: (1) free feature driving retention, (2) tiered pricing for advanced features, (3) premium tier for cross-account governance. Each needs a 2-year revenue projection.

## **Reasoning Log — Round 1**

Evaluated prd-v1.md against all 10 critique dimensions. The PRD is strongest on market timing (Why Now, score 4) and competitive differentiation (Moat, score 4). It is weakest on technical evidence (no prototype/spike data, score 2) and specificity of failure scenarios (score 2).

7/10 pass. The 3 failures are addressable but require external input: Q7 needs a technical spike (can't be solved by rewriting the PRD), Q9 needs more rigorous risk framing (solvable in a PRD revision), Q10 needs revenue modeling (requires Finance partnership).

Given that 2 of 3 failures require input beyond the PRD writer's control, the recommendation is to **ADVANCE WITH FLAGS** rather than block the pipeline. The flags are recorded in pipeline-state.md for human review before the Designer stage begins.