

Gandalf Evaluation: Quick Desktop AI Control Tower

Verdict: PASSED

Score: 12/12 questions passed | Rounds used: 1/3

Detailed Scores

#	Dimension	Rubric (1-5)	Evidence (0/1)	Pass?	Notes
1	TAM & Market	4	1	PASS	Gartner \$492M market at 45% CAGR cited with tier. Bottoms-up SAM (\$200M) and SOM (\$12M) calculated with assumptions stated.
2	Why Now	5	1	PASS	Multiple convergent signals: CSA 82% unknown agents (Apr 2026), ServiceNow Knowledge 2026 expansion (May 2026), Gartner market validation (Feb 2026), 54% C-suite fracturing (Writer 2026). Quick Desktop preview launch (Apr 2026) creates window.
3	Customer Problem Depth	4	1	PASS	Three named scenarios: Lisa's 2-week CEO investigation, David's \$180K unattributed AI spend, CISO's audit evidence gap. Grounded in CSA (82%), Logicalis (76% CIO concern), Writer (29% sabotage). Not from direct customer interviews (flagged as Open Question #3) but market-level evidence is triangulated across 5+ sources.

#	Dimension	Rubric (1-5)	Evidence (0/1)	Pass?	Notes
4	North Star Metric	5	1	PASS	"% AI tools discovered and governed" with clear rationale. Alternatives considered and rejected (MAU, cost savings, discovery count) with specific reasons. Anti-metrics defined (override rate, time-to-value). Phase gates quantified.
5	Competitive Moat	4	1	PASS	Desktop-level endpoint visibility identified as structural advantage — ServiceNow lacks desktop presence (SaaS layer), Datadog requires code instrumentation, Dynatrace needs ServiceNow for enforcement. 50+ existing integrations and knowledge graph are 18+ month head starts. Risk acknowledged: Microsoft could add governance to Copilot (Q12 FAQ).
6	Scope Discipline	5	1	PASS	Dual-scope boundary table with 14 capabilities phased across Eng v1, Proto v1, v2, v3. Specific cut rationales for each: NLP content analysis cut from Eng v1 (false positive risk), compliance reports cut (dependent on maturity framework), kill switches in v3 (safety-critical testing). Proto v1 includes placeholder pages for v2/v3 features.

#	Dimension	Rubric (1-5)	Evidence (0/1)	Pass?	Notes
7	Technical Feasibility	4	1	PASS	Hardest challenges identified: (1) privacy-preserving desktop detection — mitigated with curated allowlist approach (high precision), (2) cross-vendor API access — mitigated with top-3 vendor focus in v1, (3) governance nudge latency — Open Question #4 assigned to engineering with Week 3 deadline. Dependencies mapped with risk ratings.
8	Cannibalization Risk	4	1	PASS	Explicitly addressed in FAQ Q19: complements Bedrock Guardrails (API layer) vs. AI Control Tower (endpoint layer). Different control points, different audiences. Control Tower surfaces Guardrail status in its dashboard — strengthens both products. No cannibalization of CloudWatch or Cost Explorer.
9	Failure Mode	5	1	PASS	Three specific, falsifiable failure scenarios: (1) "spying" perception kills adoption — mitigation via transparency, opt-in, aggregation; (2) detection accuracy <80% — mitigation via allowlist approach; (3) nudge override >50% — mitigation via frequency caps. Each has quantified pivot triggers. Risk table has 6 entries with likelihood/impact /mitigation/owner.

#	Dimension	Rubric (1-5)	Evidence (0/1)	Pass?	Notes
10	Pricing & Business Model	4	1	PASS	Clear tier strategy: Discovery in Plus (teaser), full governance in Professional/Enterprise. Land-and-expand motion: free → Plus → Professional. Revenue path via tier upgrades. V2 pricing gate at 50K users. FAQ Q18 details the model. Strategic retention value articulated (stickier Quick Desktop for enterprise).
11	Solution Direction Deliberation	5	1	PASS	Solution Lineage table exists with 6 rows. Each traces to a specific Opportunity from the research tree. Rejected alternatives named with specific rationales: B/Integration-only rejected for missing shadow AI; C/Notification nudges rejected for weakest enforcement; B/Activity-based estimation rejected for accuracy concerns. No "out of scope" hand-waving — each rejection cites a concrete cost.

#	Dimension	Rubric (1-5)	Evidence (0/1)	Pass?	Notes
12	Product Completeness	4	1	PASS	Proto v1 scope covers 14 capabilities including placeholder pages for v2/v3 features (integration marketplace, compliance reports, enforcement page, template library). Dual-scope table explicitly maps Eng v1 vs Proto v1. Research Interaction Pattern Benchmarking cited ServiceNow's 5-pillar navigation and Datadog's 6-section sidebar — Proto v1 covers all 5 dimensions (Discover, Observe, Govern, Secure, Measure) plus Admin Console and Integration Marketplace = 7 navigation sections, meeting or exceeding competitor surface.

Questions That Passed

- TAM & Market:** Gartner-backed \$492M market with bottoms-up SAM/SOM calculation showing assumptions.
- Why Now:** Five convergent timing signals from Q4 2025 — May 2026, each independently sourced.
- Customer Problem Depth:** Three named personas with day-in-the-life narratives; 5 enumerated workarounds with quantified costs; market evidence triangulated across CSA, Logicalis, Writer, EY, Gartner.
- North Star Metric:** Well-justified metric choice with rejected alternatives, anti-metrics, and quantified phase gates.
- Competitive Moat:** Desktop endpoint visibility is a genuine structural advantage; 18+ month head start on knowledge graph; pricing moat (\$20 vs enterprise ITSM).
- Scope Discipline:** Dual-scope table with 14 phased capabilities; specific cut rationales; Proto v1 includes vision pages for cut features.
- Technical Feasibility:** Hardest challenges named with specific mitigations; open questions assigned with deadlines.
- Cannibalization Risk:** Complementary positioning clearly articulated — different control points (API vs. endpoint).
- Failure Mode:** Three specific failure scenarios with quantified pivot triggers and owner assignments.
- Pricing & Business Model:** Tier strategy, land-and-expand motion, retention value, and v2 pricing gate defined.

11. **Solution Direction Deliberation:** Complete Solution Lineage table with named rejections and specific rationales.
12. **Product Completeness:** Proto v1 covers 7 navigation sections (vs. ServiceNow's 5 pillars + dashboard), including placeholder pages for v2/v3 features.

Questions Flagged for Human Review

None — all 12 questions passed in Round 1.

Minor Observations (not blocking)

1. **Customer voice is market-level, not Quick Desktop-specific.** All evidence comes from industry surveys (CSA, Writer, Logicalis) not from Quick Desktop pilot customers. Open Question #3 (UXR interviews by Week 4) addresses this. Recommend prioritizing these interviews before v1 design lock.
2. **M365 Admin API dependency is a real risk.** FAQ Q10 acknowledges this but the mitigation ("ship without Copilot cost data") weakens the cross-vendor cost attribution value proposition. Recommend parallel-tracking the Microsoft partnership (Open Question #2) with a fallback cost estimation approach.
3. **Privacy framework is critical path.** Legal review (6-8 week cycle per Dependencies Map) could delay launch. Recommend starting legal engagement immediately, not after design phase.

Approved Changes to PRD

No changes required. PRD v1 passes Gandalf Gate as-is.

Reasoning Log

Round 1

Evaluated PRD v1 (Quick Desktop AI Control Tower) against all 12 Gandalf questions. The PRD demonstrates strong strategic framing with the "cockpit, not control tower" positioning, evidence-backed claims across 18 sources, and a well-structured dual-scope boundary that addresses the product completeness concern proactively. The Solution Lineage table is thorough with specific rejection rationales. North Star metric selection is well-justified with rejected alternatives. Failure modes are specific and falsifiable with quantified pivot triggers. All 12 questions pass. No Round 2 needed.